

MapReduce

MapReduce

- MapReduce [OSDI'04] provides
 - Automatic parallelization, distribution
 - I/O scheduling
 - Load balancing
 - Network and data transfer optimization
 - Fault tolerance
 - Handling of machine failures
- Need more power: Scale out, not up!
 - Large number of commodity servers as opposed to some high end specialized servers


Apache Hadoop:

Open source
implementation of
MapReduce

MapReduce workflow

Input Data

Output Data



Map

extract something you
care about from each
record

Reduce

aggregate,
summarize, filter,
or transform

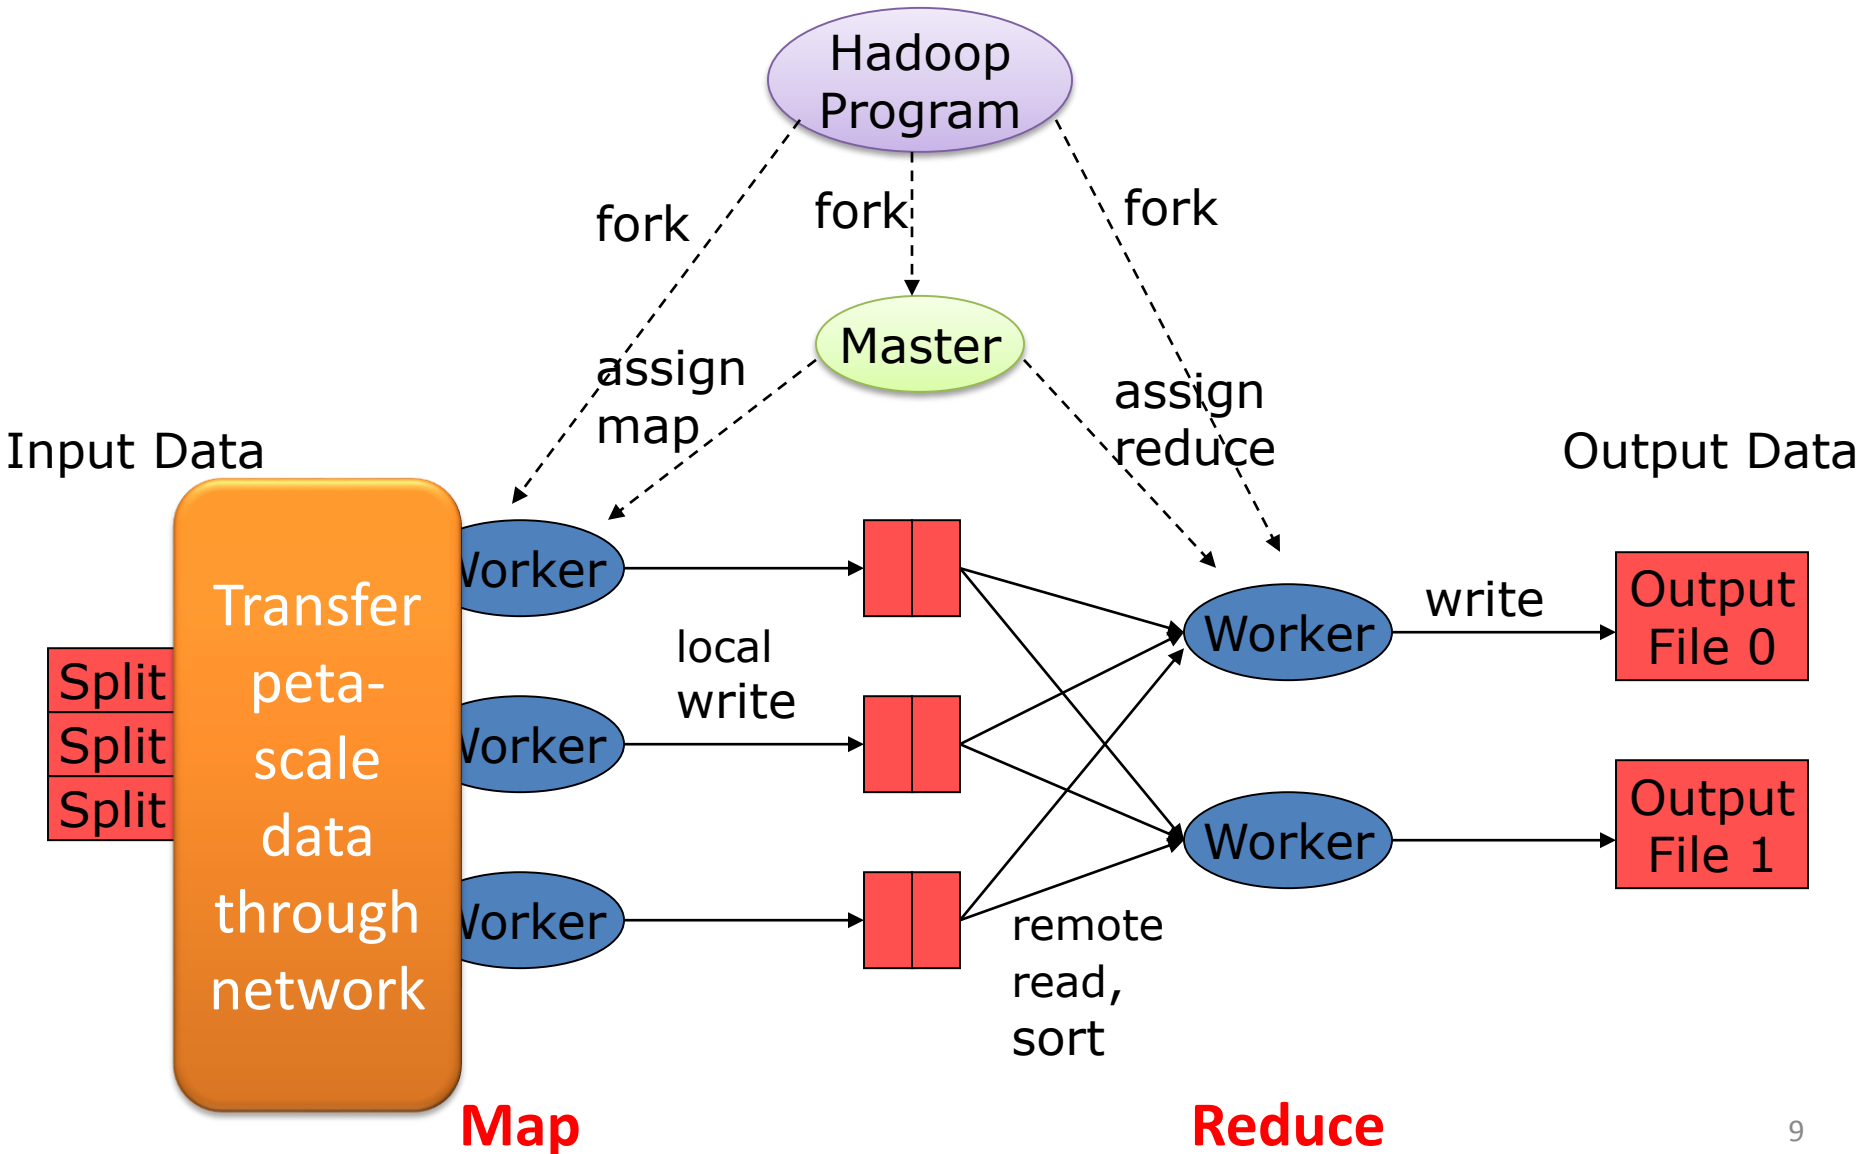
Example: Word Count

Input Files

Apple Orange Mango
Orange Grapes Plum

Apple Plum Mango
Apple Apple Plum


MapReduce



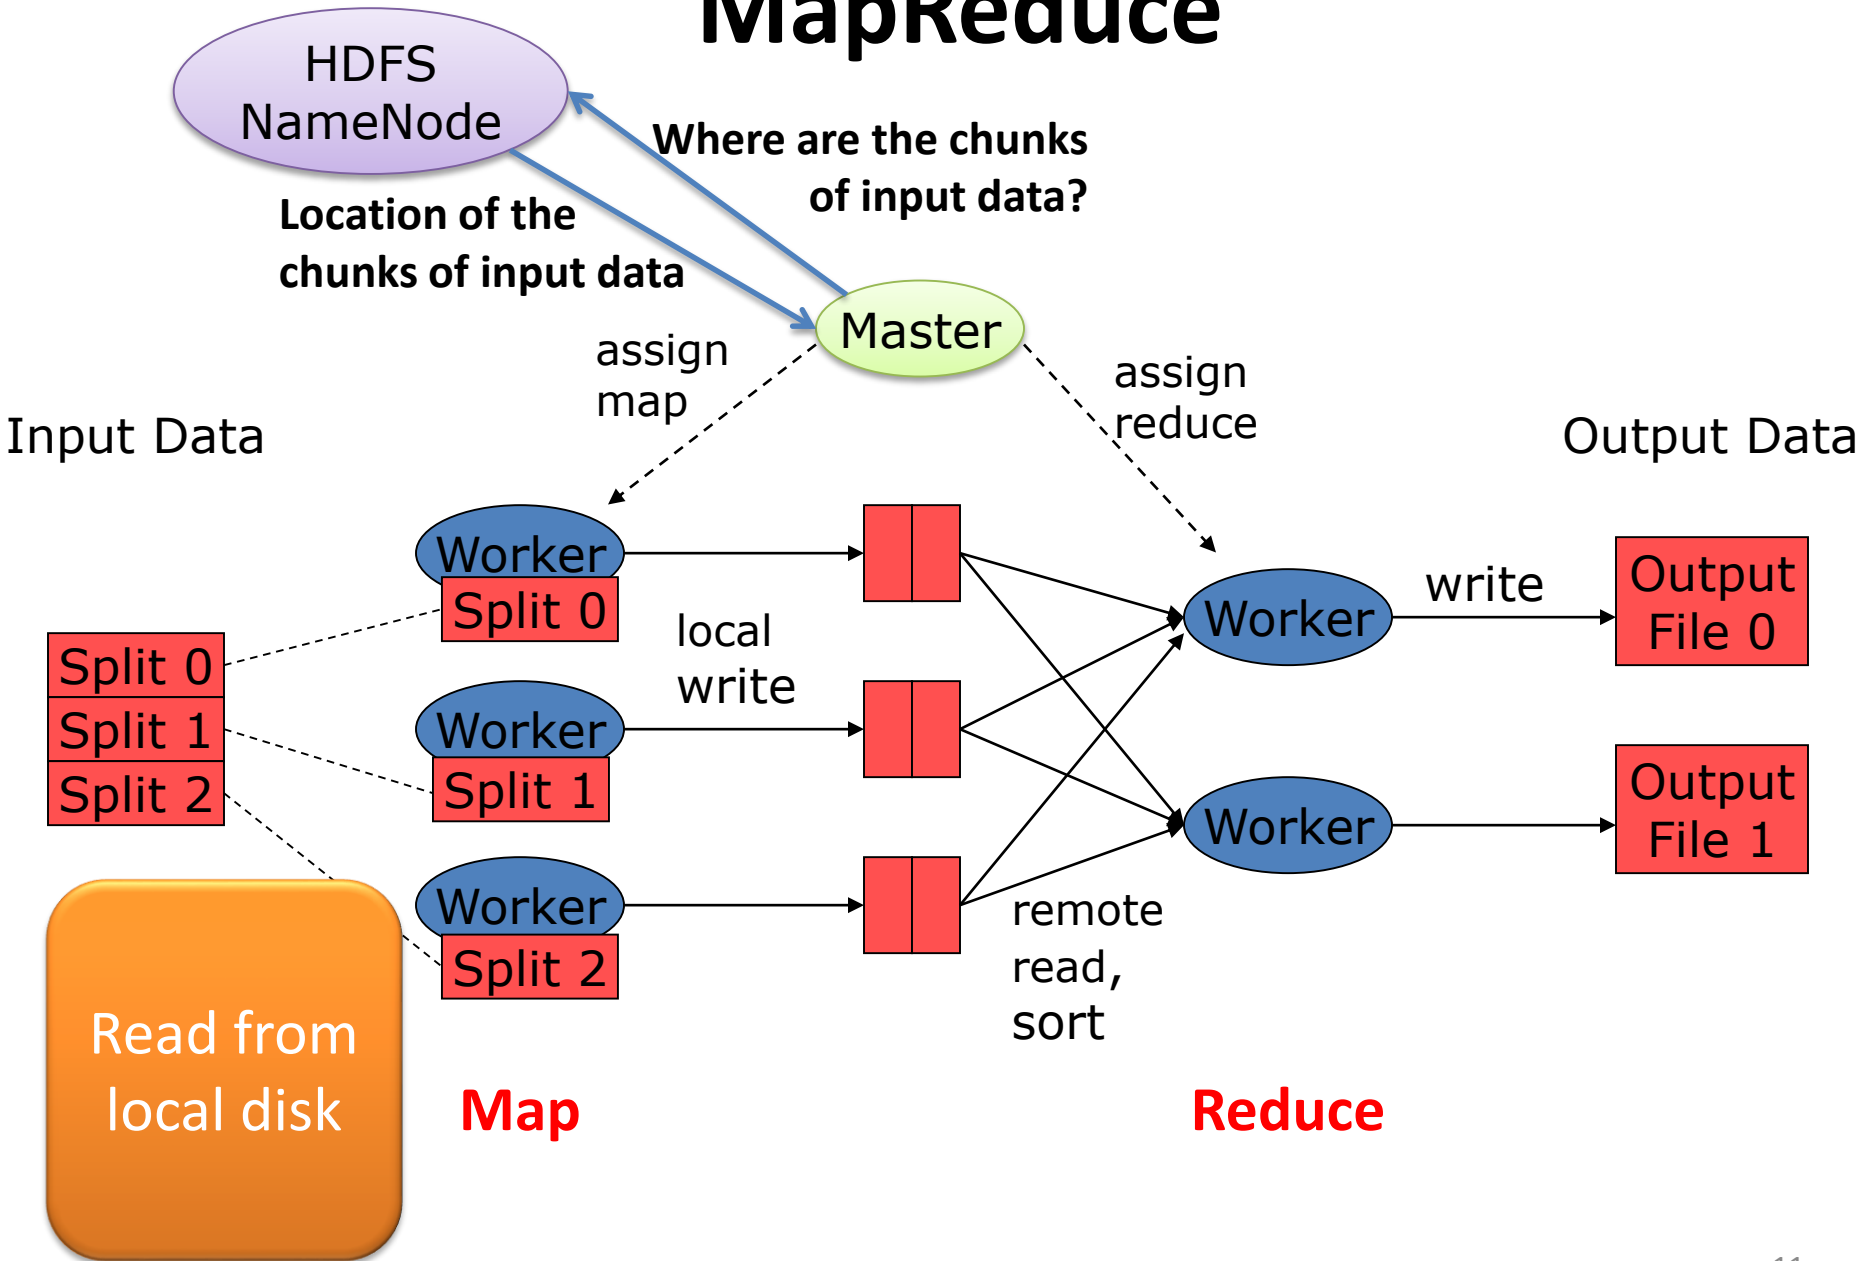
Google File System (GFS)

Hadoop Distributed File System (HDFS)

- Split data and store 3 replica on commodity servers



MapReduce



Failure in MapReduce

- **Failures** are **norm** in commodity hardware
- **Worker** failure
 - Detect failure via periodic **heartbeats**
 - **Re-execute** in-progress map/reduce tasks
- **Master** failure
 - Single point of failure; Resume from Execution Log
- **Robust**
 - Google's experience: **lost 1600 of 1800 machines once!**, but **finished fine**


```
public class WordCount {
```

```
    public static class Map extends Mapper<LongWritable, Text, Text, IntWritable> {  
        private final static IntWritable one = new IntWritable(1);  
        private Text word = new Text();
```

Mapper

```
        public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {  
            String line = value.toString();  
            StringTokenizer tokenizer = new StringTokenizer(line);  
            while (tokenizer.hasMoreTokens()) {  
                word.set(tokenizer.nextToken());  
                context.write(word, one);  
            }  
        }  
    }  
}
```

```
    public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {
```

```
        public void reduce(Text key, Iterable<IntWritable> values, Context context)  
            throws IOException, InterruptedException {  
            int sum = 0;  
            for (IntWritable val : values) {  
                sum += val.get();  
            }  
            context.write(key, new IntWritable(sum));  
        }  
    }  
}
```

Reducer

```
    public static void main(String[] args) throws Exception {
```

```
        Configuration conf = new Configuration();
```

```
        Job job = new Job(conf, "wordcount");
```

```
        job.setOutputKeyClass(Text.class);
```

```
        job.setOutputValueClass(IntWritable.class);
```

```
        job.setMapperClass(Map.class);
```

```
        job.setReducerClass(Reduce.class);
```

```
        job.setInputFormatClass(TextInputFormat.class);
```

```
        job.setOutputFormatClass(TextOutputFormat.class);
```

```
        FileInputFormat.addInputPath(job, new Path(args[0]));
```

```
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
```

```
        job.waitForCompletion(true);
```

```
    }  
}
```

Run this program as
a MapReduce job

Contents

- **Motivation**
- Design overview
 - Write Example
 - Record Append
- Fault Tolerance & Replica Management
- Conclusions

Motivation: Large Scale Data Storage


- Manipulate large (**Peta Scale**) sets of data
- Large number of machine with **commodity hardware**
- Component failure is the norm
- Goal: **Scalable, high performance, fault tolerant**
distributed file system

Why a new file system?

- None designed for their failure model
- Few scale as highly or dynamically and easily
- Lack of special primitives for large distributed computation

What should expect from GFS

- Designed for Google's application
 - Control of both file system and application
 - Applications use a few specific access patterns
 - Append to large files
 - Large streaming reads
 - **Not** a good fit for
 - low-latency data access
 - lots of small files, multiple writers, arbitrary file modifications
- Not POSIX, although mostly traditional
 - Specific operations: RecordAppend




Contents

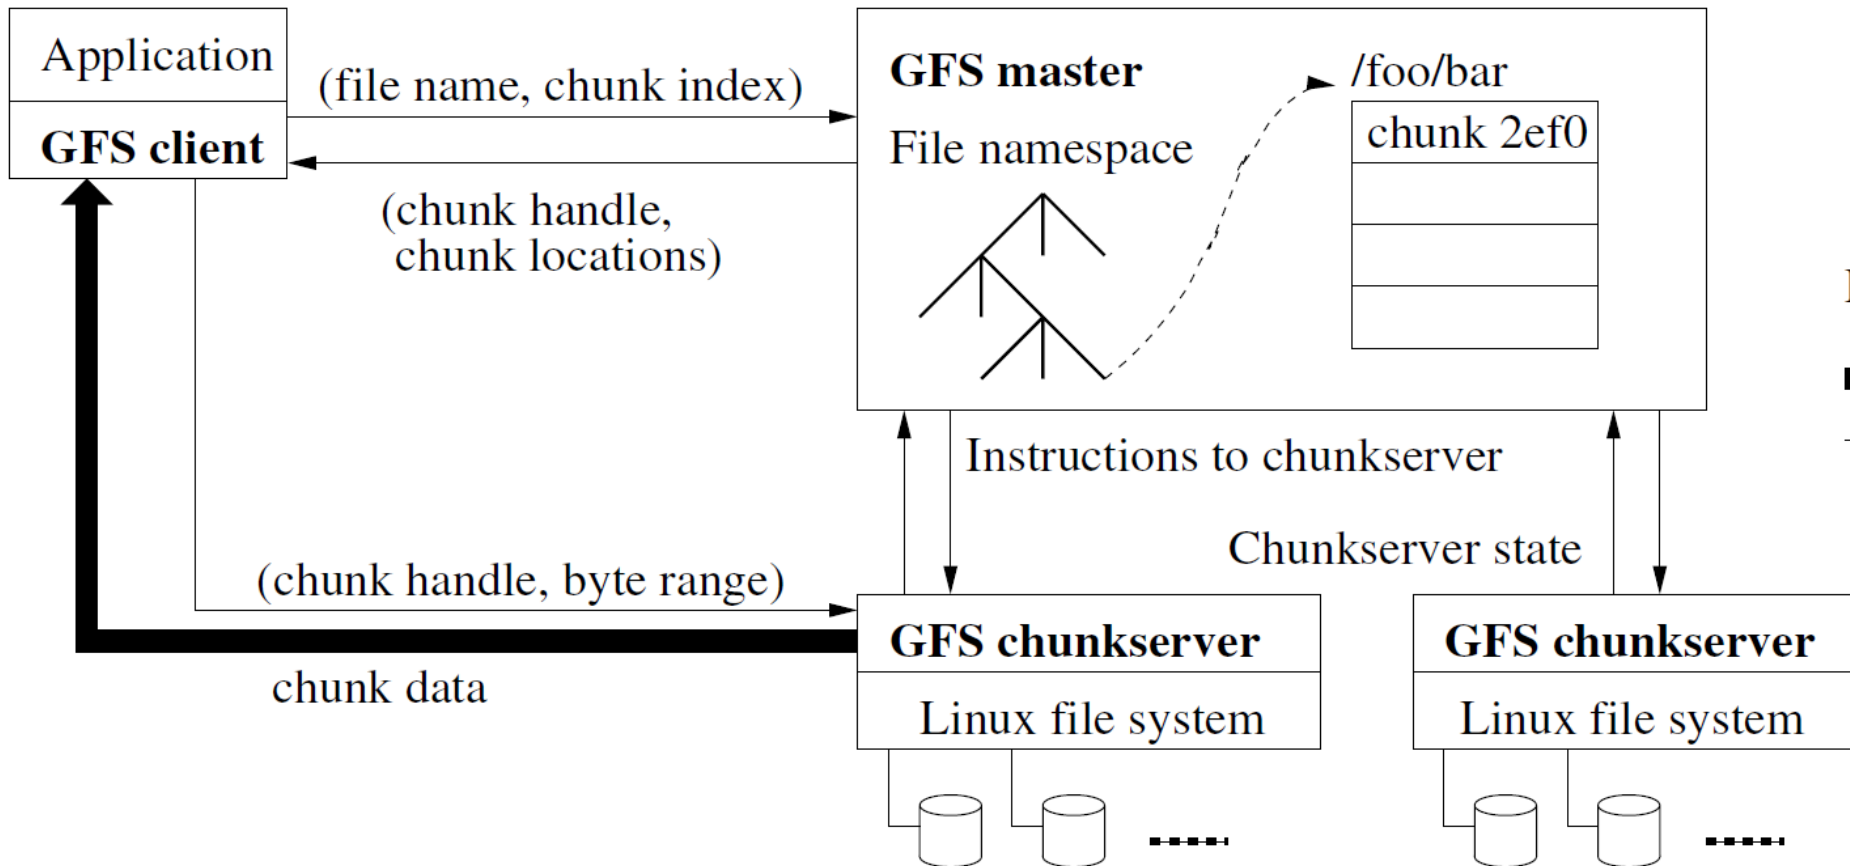
- Motivation
- **Design overview**
 - Write Example
 - Record Append
- Fault Tolerance & Replica Management
- Conclusions

Components

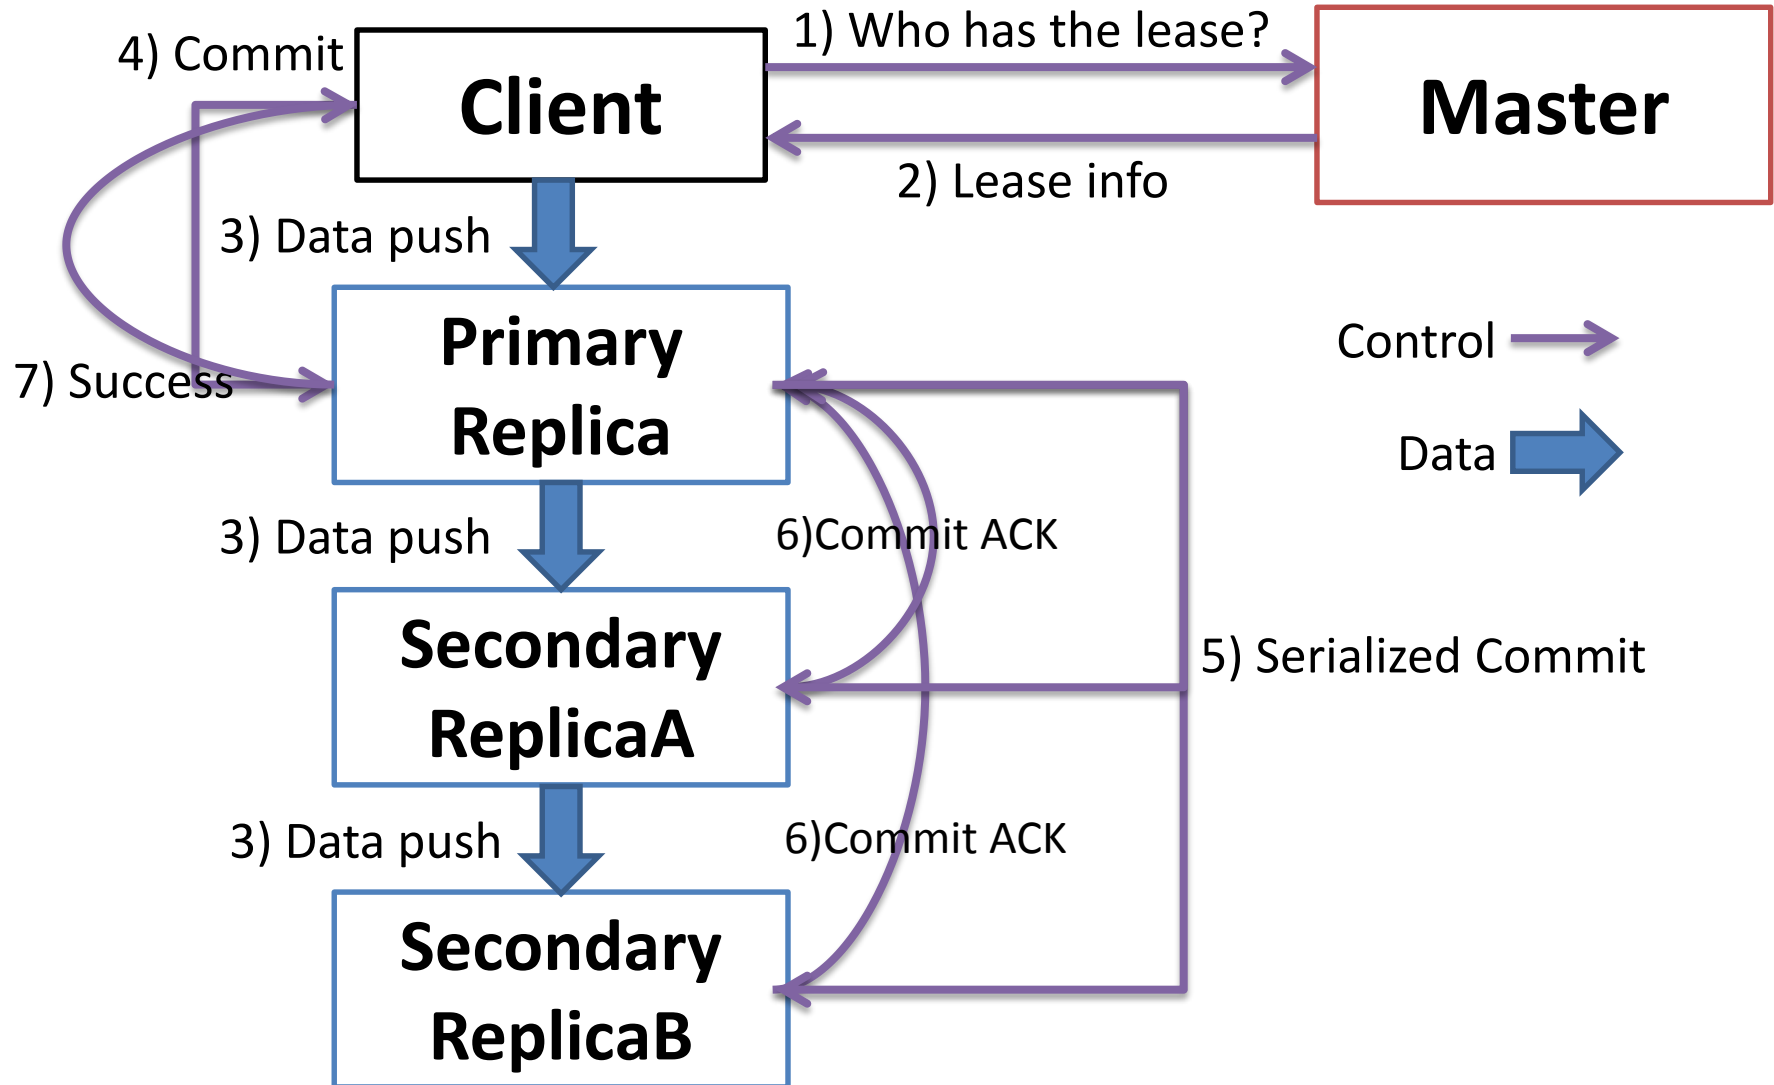
- **Master (NameNode)**
 - Manages metadata (namespace)
 - Not involved in data transfer
 - Controls allocation, placement, replication
- **Chunkserver (DataNode)**
 - Stores chunks of data
 - No knowledge of GFS file system structure
 - Built on local linux file system



GFS Architecture



Write(filename, offset, data)



RecordAppend(filename, data)

- Significant use in distributed apps. For example at Google production cluster:
 - 21% of bytes written
 - 28% of write operations
- **Guaranteed:** All data appended at least once as a single consecutive byte range
- Same basic structure as write
 - Client obtains information from master
 - Client sends data to data nodes (chunkservers)
 - Client sends “append-commit”
 - Lease holder serializes append
- **Advantage:** Large number of concurrent writers with minimal coordination

RecordAppend (2)

- Record size is limited by chunk size
- When a record does not fit into available space,
 - chunk is padded to end
 - and client retries request.

Contents

- Motivation
- Design overview
 - Write Example
 - Record Append
- **Fault Tolerance & Replica Management**
- Conclusions

Fault tolerance

- Replication
 - High availability for reads
 - User controllable, default 3 (non-RAID)
 - Provides read/seek bandwidth
 - Master is responsible for directing re-replication if a data node dies
- Online checksumming in data nodes
 - Verified on reads

Replica Management

- Bias towards **topological** spreading
 - Rack, data center
- **Rebalancing**
 - Move chunks around to balance disk fullness
 - Gently fixes imbalances due to:
 - Adding/removing data nodes

Replica Management (Cloning)

- Chunk replica lost or corrupt
- **Goal:** minimize app disruption and data loss
 - Approximately in priority order
 - More replica missing-> priority boost
 - Deleted file-> priority decrease
 - Client blocking on a write-> large priority boost
 - Master directs copying of data
- Performance on a production cluster
 - Single failure, full recovery (600GB): 23.2 min
 - Double failure, restored 2x replication: 2min

Garbage Collection

- Master does **not** need to have a **strong knowledge** of what is stored on each data node
 - Master regularly scans namespace
 - After GC interval, deleted files are removed from the namespace
 - Data node periodically polls Master about each chunk it knows of.
 - If a chunk is forgotten, the master tells data node to delete it.

Limitations

- Master is a central point of failure
- Master can be a scalability bottleneck
- Latency when opening/stating thousands of files
- Security model is weak

Conclusion

- Inexpensive commodity components can be the basis of a large scale reliable system
- Adjusting the API, e.g. RecordAppend, can enable large distributed apps
- Fault tolerant
- Useful for many similar apps